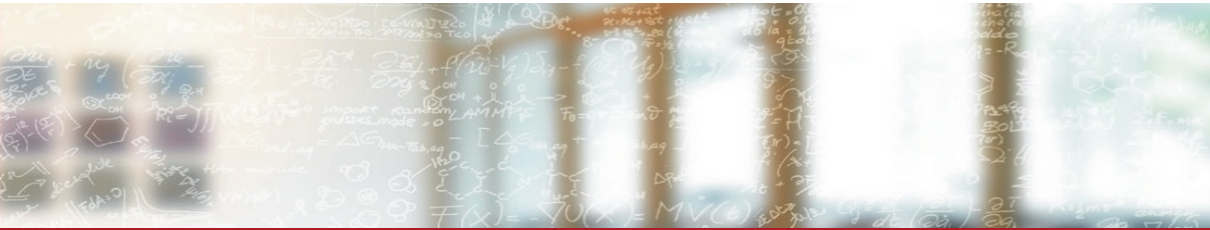




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Introduction to the Piz Daint environment

CSCS-USI Summer School 2019

*Vasileios Karakasis, CSCS*

July 15, 2019

# Overview

- Accessing CSCS
- Compiling my code
- Running my code
- Editing my code
- Transferring files from/to CSCS
- Repository of the course

# Piz Daint

## Computing nodes

Piz Daint is a hybrid cluster of Cray XC40/XC50 nodes

- Hybrid nodes (XC50)
  - 5320 total
  - Intel® Xeon® E5-2690 v3 @ 2.60GHz (12 cores, 64GB RAM, Haswell)
  - NVIDIA® Tesla® P100 16GB (Pascal)
- Multicore nodes
  - 1813 nodes
  - Two Intel® Xeon® E5-2695 v4 @ 2.10GHz (2 x 18 cores, 64/128 GB RAM, Broadwell)
- Login nodes
  - 5 total
  - Intel® Xeon® CPU E5-2650 v3 @ 2.30GHz (10 cores, 256 GB RAM, Haswell)
- Aries routing and communications ASIC, and Dragonfly network topology

# Piz Daint

## Filesystems

- `/scratch`: High performance Lustre filesystem accessible from the computing nodes
  - Environment variable `$SCRATCH` points to it
  - Total capacity: 6.2 PB
  - Must be used for heavy I/O
- `/users`: GPFS filesystem for the users' homes
- `/project`, `/store`: Long-term storage for computational projects

More on [https://user.cscs.ch/storage/file\\_systems/](https://user.cscs.ch/storage/file_systems/)

# Accessing Piz Daint

- Accessible through SSH
- Piz Daint is not directly accessible from the outside world:
  - `ela` → `daint10x` → `nidxxxxxx`

Two-steps process:

1. Login to the frontend, forwarding X11 (will be needed the second day)
2. Move to the login nodes of Piz Daint

```
# Login to the frontend first  
ssh -Y courseXX@ela.cscs.ch  
ssh -Y daint
```

# Programming Environments

## Cray Linux Programming Environment

- 4 compilers available: CCE, GNU, INTEL, PGI
- 4 predefined Programming Environments:
  - PrgEnv-cray (default), PrgEnv-gnu, PrgEnv-intel, PrgEnv-pgi
  - `echo $PE_ENV` to get the current PrgEnv
- 3 wrappers available: `ftn` (Fortran), `cc` (C), `CC` (C++)
  - Required for compiling MPI programs
  - They set appropriate optimisation flags for the target architecture (CPU or GPU)
  - They provide a sort of portability across the programming environments

# Managing programming environments

Daint uses *Environment Modules* (TMod) for managing the programming environments and the software packages:

- Dynamic modification of a user's environment via *modulefiles*.
- All programming environments and software on Daint is available through modules.
- The compiler wrappers will detect the loaded programming environment and automatically set the correct flags and libraries.

# Managing programming environments

## Listing modules

– `module list`

```
3. → karakasv@ela1:~ [Default]
stud33@daint105:~> module list
Currently Loaded Modulefiles:
 1) modules/3.2.10.6
 2) cce/8.7.3
 3) craype-network-aries
 4) craype/2.5.15
 5) cray-libsci/18.07.1
 6) udreg/2.3.2-6.0.7.1_5.13__g5196236.ari
 7) ugni/6.0.14.0-6.0.7.1_3.13__gea11d3d.ari
 8) pmi/5.0.14
 9) dmapp/7.1.1-6.0.7.1_5.45__g5a674e0.ari
10) gni-headers/5.0.12.0-6.0.7.1_3.11__g3b1768f.ari
11) xpmem/2.2.15-6.0.7.1_5.11__g7549d06.ari
12) job/2.2.3-6.0.7.1_5.43__g6c4e934.ari
13) dvs/2.7_2.2.118-6.0.7.1_10.2__g58b37a2
14) alps/6.6.43-6.0.7.1_5.45__ga796da32.ari
15) rca/2.2.18-6.0.7.1_5.47__g2aa4f39.ari
16) atp/2.1.2
17) perftools-base/7.0.3
18) PrgEnv-cray/6.0.4
19) cray-mpich/7.7.2
20) slurm/17.11.12.cscs-1
21) craype-haswell
22) xalt/daint-2016.11
stud33@daint105:~> █
```



# Managing programming environments

## Switching programming environments

- Switch to the PGI programming environment
- module switch

```
3. -> karakasv@ela1:~ [Default]
stud33@daint107:~> module switch PrgEnv-cray/6.0.4 PrgEnv-pgi
stud33@daint107:~> module list
Currently Loaded Modulefiles:
  1) modules/3.2.10.6
  2) pgi/18.5.0
  3) craype-haswell
  4) craype-network-aries
  5) craype/2.5.15
  6) cray-mpich/7.7.2
  7) slurm/17.11.12.cscs-1
  8) xalt/daint-2016.11
  9) udreg/2.3.2-6.0.7.1_5.13__g5196236.ari
 10) ugni/6.0.14.0-6.0.7.1_3.13__gea11d3d.ari
 11) pmi/5.0.14
 12) dmapp/7.1.1-6.0.7.1_5.45__g5a674e0.ari
 13) gni-headers/5.0.12.0-6.0.7.1_3.11__g3b1768f.ari
 14) xpmem/2.2.15-6.0.7.1_5.11__g7549d06.ari
 15) job/2.2.3-6.0.7.1_5.43__g6c4e934.ari
 16) dvs/2.7.2.2.118-6.0.7.1_10.2__g58b37a2
 17) alps/6.6.43-6.0.7.1_5.45__ga796da32.ari
 18) rca/2.2.18-6.0.7.1_5.47__g2aa4f39.ari
 19) atp/2.1.2
 20) perftools-base/7.0.3
 21) PrgEnv-pgi/6.0.4
stud33@daint107:~> ftn -V

pgf90 18.5-0 64-bit target on x86-64 Linux -tp haswell-64
PGI Compilers and Tools
Copyright (c) 2018, NVIDIA CORPORATION. All rights reserved.
stud33@daint107:~> █
```

# Managing programming environments

## Switching back to the Cray programming environment

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> module switch PrgEnv-pgi/6.0.4 PrgEnv-cray
stud33@daint107:~> module list
Currently Loaded Modulefiles:
  1) modules/3.2.10.6
  2) slurm/17.11.12.cscs-1
  3) xalt/daint-2016.11
  4) cce/8.7.3
  5) craype-haswell
  6) craype-network-aries
  7) craype/2.5.15
  8) cray-mpich/7.7.2
  9) cray-libsci/18.07.1
 10) udreg/2.3.2-6.0.7.1_5.13__g5196236.ari
 11) ugni/6.0.14.0-6.0.7.1_3.13__gea11d3d.ari
 12) pmi/5.0.14
 13) dmapp/7.1.1-6.0.7.1_5.45__g5a674e0.ari
 14) gni-headers/5.0.12.0-6.0.7.1_3.11__g3b1768f.ari
 15) xpmem/2.2.15-6.0.7.1_5.11__g7549d06.ari
 16) job/2.2.3-6.0.7.1_5.43__g6c4e934.ari
 17) dvs/2.7_2.2.118-6.0.7.1_10.2__g58b37a2
 18) alps/6.6.43-6.0.7.1_5.45__ga796da32.ari
 19) rca/2.2.18-6.0.7.1_5.47__g2aa4f39.ari
 20) atp/2.1.2
 21) perftools-base/7.0.3
 22) PrgEnv-cray/6.0.4
stud33@daint107:~> ftn -V
Cray Fortran : Version 8.7.3  Wed Jul 10, 2019  17:47:52
stud33@daint107:~> █
```

# Managing programming environments

## Loading and unloading modules

- `module load [MODULE_NAME]`
- `module unload [MODULE_NAME]`

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> module load cray-hdf5
stud33@daint107:~> which h5dump
/opt/cray/pe/hdf5/1.10.2.0/bin/h5dump
stud33@daint107:~> module unload cray-hdf5
stud33@daint107:~> which h5dump
which: no h5dump in (/opt/cray/pe/perftools/7.0.3/bin:/opt/cray/pe/papi/5.6.0.3/bin:/opt/cray/rca/2.2.18-6.0.7.1_5.47__g2aa4f3
9.ari/bin:/opt/cray/alps/6.6.43-6.0.7.1_5.45__ga796da32.ari/sbin:/opt/cray/job/2.2.3-6.0.7.1_5.43__g6c4e934.ari/bin:/opt/cray/
pe/mpt/7.7.2/gni/bin:/opt/cray/pe/craype/2.5.15/bin:/opt/cray/pe/cce/8.7.3/binutils/x86_64/x86_64-pc-linux-gnu/bin:/opt/cray/p
e/cce/8.7.3/binutils/cross/x86_64-aarch64/aarch64-linux-gnu/..bin:/opt/cray/pe/cce/8.7.3/Utils/x86_64/bin:/apps/daint/UES/xal
t/0.7.6/bin:/opt/slurm/17.11.12.cscs/bin:/opt/cray/pe/modules/3.2.10.6/bin:/opt/slurm/default/bin:/apps/daint/system/bin:/apps
/common/system/bin:/users/stud33/bin:/usr/local/bin:/usr/bin:/bin:/usr/bin/X11:/usr/lib/mit/bin:/usr/lib/mit/sbin:/opt/cray/pe
/bin)
stud33@daint107:~> h5dump
If 'h5dump' is not a typo you can use command-not-found to lookup the package that contains it, like this:
  cnf h5dump
stud33@daint107:~> █
```

# Managing programming environments

## Checking available modules

- The daint-gpu makes available the CSCS software stack built for the hybrid nodes of the system

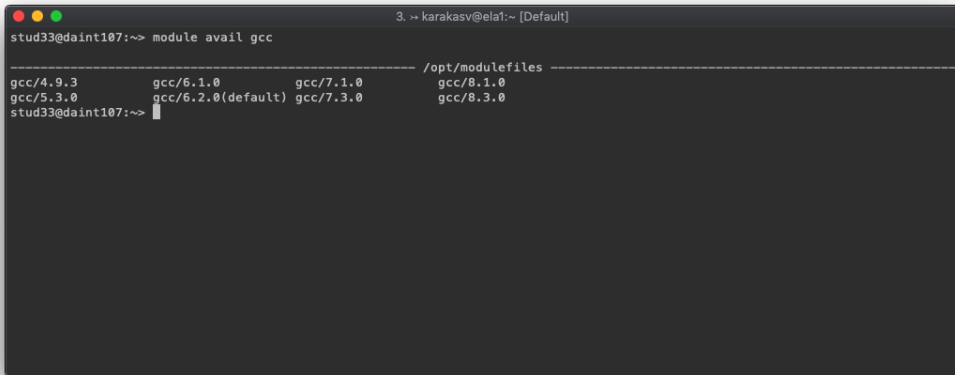
```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> module load daint-gpu
stud33@daint107:~> module avail

----- /apps/daint/UES/jenkins/6.0.UP07/gpu/easybuild/tools/modules/all -----
CubeGUI/4.4(default)
CubeLib/4.4
CubeLib/4.4-CrayCCE-18.08
CubeLib/4.4-CrayGNU-18.08(default)
CubeLib/4.4-CrayIntel-18.08
CubeLib/4.4-CrayPGI-18.08
CubeW/4.4-CrayCCE-18.08
CubeW/4.4-CrayGNU-18.08(default)
CubeW/4.4-CrayIntel-18.08
CubeW/4.4-CrayPGI-18.08
Dimemas/5.3.6(default)
Dimemas/latest
Extrae/3.5.4-CrayGNU-18.08(default)
IPM/2.0.6-CrayCCE-18.08
IPM/2.0.6-CrayGNU-18.08(default)
IPM/2.0.6-CrayIntel-18.08
IPM/2.0.6-CrayPGI-18.08
MUST/v1.6-rc1-CrayGNU-18.08(default)
Paraver/4.7.2(default)
callgraphplugin/0.1(default)
clusteringsuite/2.6.7
clusteringsuite/2.6.8(default)
clusteringsuite/latest
ddt/18.2.1-Suse-12(default)
ddt/18.3-Suse-12
ddt/19.0-Suse-12
folding/1.3.2(default)
folding/latest
go/1.12.1.linux-amd64(default)
gperftools/2.7(default)
gprof2dot/2017.9.19-CrayGNU-18.08-python3(default)
graphviz/2.40.1(default)
inspector/2018(default)
inspector/2019
inspector/2019_update4
jengafettplugin/0.1(default)
libelf/0.8.13(default)
likwid/4.3.3-perf_event(default)
```

# Managing programming environments

## Checking available modules

- Check available versions of a software



```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> module avail gcc

----- /opt/modulefiles -----
gcc/4.9.3      gcc/6.1.0      gcc/7.1.0      gcc/8.1.0
gcc/5.3.0      gcc/6.2.0(default) gcc/7.3.0      gcc/8.3.0
stud33@daint107:~> 
```

# Managing programming environments

Get information about a module

- Environment variables set, paths etc.

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> module show gcc
-----
/opt/modulefiles/gcc/6.2.0:

conflict      gcc
conflict      gcc-cross-aarch64
prepend-path  PATH /opt/gcc/6.2.0/bin
prepend-path  MANPATH /opt/gcc/6.2.0/snos/share/man
prepend-path  INFOPATH /opt/gcc/6.2.0/snos/share/info
prepend-path  LD_LIBRARY_PATH /opt/gcc/6.2.0/snos/lib64
setenv        GCC_PATH /opt/gcc/6.2.0
setenv        GCC_VERSION 6.2.0
setenv        GNU_VERSION 6.2.0
-----

stud33@daint107:~> █
```

# Managing programming environments

Get help for a module

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> module help CP2K

----- Module Specific Help for 'CP2K/6.1-CrayGNU-18.08-cuda-9.1' -----

Description
=====
CP2K is a freely available (GPL) program, written in Fortran 95, to perform atomistic and molecular
simulations of solid state, liquid, molecular and biological systems. It provides a general framework for different
methods such as e.g. density functional theory (DFT) using a mixed Gaussian and plane waves approach (GPW), and
classical pair and many-body potentials.

More information
=====
- Homepage: http://www.cp2k.org/

stud33@daint107:~> █
```

# Running on Piz Daint

The job scheduler

Piz Daint uses native SLURM for running jobs on the compute nodes. There are three ways of submitting a job:

1. Interactively from the login nodes using the `srun` command.
2. By submitting a job script using the `sbatch` command.
3. By pre-allocating resources using the `salloc` command.



# Running on Piz Daint

Using the `srun` command

Necessary and useful options:

- `-C gpu`: requests allocation on the hybrid (GPU) nodes (required)
- `--reservation=summer_school`
  - 64 nodes valid until Jul. 25 @ 13:00.
- `-N 2`: number of compute nodes (default is 1)
- `-n 2`: number of MPI tasks (default is 1)
- `-t 5`: maximum duration of the job (default is 30min)
  - Allows to get an allocation quicker
  - Job will be killed if time limit is reached
  - Maximum time slot for a job is 24h

More on <https://user.cscs.ch/access/running>

# Running on Piz Daint

Using the `srun` command

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> srun --reservation=summer_school_test -Cgpu -t1 -N2 hostname
srun: job 15391913 queued and waiting for resources
srun: job 15391913 has been allocated resources
nid02296
nid02822
stud33@daint107:~> srun --reservation=summer_school_test -Cgpu -t1 -n2 hostname
srun: job 15391932 queued and waiting for resources
srun: job 15391932 has been allocated resources
nid02296
nid02296
stud33@daint107:~> █
```

# Running on Piz Daint

Using the sbatch command

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> cat job.sh
#!/bin/bash
#SBATCH -J 'my_first_job'
#SBATCH -C gpu
#SBATCH -N 2
#SBATCH --reservation=summer_school_test
#SBATCH -o myjob.out
#SBATCH -e myjob.err

echo "My job ID is $SLURM_JOB_ID"
srun hostname
stud33@daint107:~> sbatch job.sh
Submitted batch job 15391950
stud33@daint107:~> squeue -j 15391950
  JOBID  USER ACCOUNT      NAME  ST REASON   START_TIME          TIME  TIME_LEFT NODES  CPUS
15391950 stud33  std01  my_first_job CF None    18:43:19          0:02    29:58      2   48
stud33@daint107:~> squeue -j 15391950
  JOBID  USER ACCOUNT      NAME  ST REASON   START_TIME          TIME  TIME_LEFT NODES  CPUS
15391950 stud33  std01  my_first_job R  None    18:43:19          0:09    29:51      2   48
stud33@daint107:~> █
```

# Running on Piz Daint

Using the `sbatch` command – Examining the output

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> cat myjob.err
stud33@daint107:~> cat myjob.out
My job ID is 15391950
nid02296
nid02822

Batch Job Summary Report for Job "my_first_job" (15391950) on daint
-----
Submit          Eligible          Start          End          Elapsed  Timelimit
-----
2019-07-10T18:43:08 2019-07-10T18:43:08 2019-07-10T18:43:19 2019-07-10T18:43:34 00:00:15 00:30:00
-----
Username  Account  Partition  NNodes  Energy
-----
stud33    std01    normal     2        1.44K joules

This job did not utilize any GPUs

-----
Scratch File System  Files  Quota
-----
/scratch/snx3000     33    1000000
```

# Running on Piz Daint

Using the salloc command

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> salloc --reservation=summer_school_test -Cgpu -t1 -N2
salloc: Pending job allocation 15392032
salloc: job 15392032 queued and waiting for resources
salloc: job 15392032 has been allocated resources
salloc: Granted job allocation 15392032
stud33@daint107:~> srun -N2 hostname
nid02296
nid02822
stud33@daint107:~> srun -N1 hostname
nid02296
stud33@daint107:~> srun -N4 hostname
srun: error: Only allocated 2 nodes asked for 4
stud33@daint107:~> exit
salloc: Relinquishing job allocation 15392032
stud33@daint107:~> █
```

# Running on Piz Daint

## Other useful commands

- `squeue [OPTIONS]`: Check the status of the system job queue
  - Useful options: `-u [USERNAME]`, `-j [JOBID]`
- `scancel [JOBID]`: Cancel a job
- `scontrol`: Detailed information about partitions, reservations, computing nodes etc.

# Running on Piz Daint

## Other useful commands

```
3. → karakasv@ela1:~ [Default]
stud33@daint107:~> scontrol show reservation summer_school
ReservationName=summer_school StartTime=Mon 12:00 EndTime=25 Jul 13:00 Duration=10-01:00:00
Nodes=nid[6144-6207] NodeCnt=64 CoreCnt=768 Features=gpu PartitionName=normal Flags=
TRES=cpu=1536
Users=(null) Accounts=std01,u3 Licenses=(null) State=INACTIVE BurstBuffer=(null) Watts=n/a

stud33@daint107:~> scontrol show partition normal
PartitionName=normal
AllowGroups=ALL AllowAccounts=ALL AllowQos=ALL
AllocNodes=ALL Default=YES QoS=N/A
DefaultTime=00:30:00 DisableRootJobs=NO ExclusiveUser=NO GraceTime=0 Hidden=NO
MaxNodes=2400 MaxTime=1-00:00:00 MinNodes=1 LLN=NO MaxCPUsPerNode=UNLIMITED
Nodes=nid[00004-00007,00012-00024,00026-00062,00064-00067,00072-00126,00128-00190,00192-00195,00200-00254,00260-00318,00320-00323,00328-00382,00388-00446,00456-00510,00516-00574,00576-00579,00584-00638,00644-00702,00704-00707,00712-00766,00772-00830,00832-00835,00840-00894,00900-00958,00960-00963,00968-01022,01028-01086,01088-01150,01152-01192,01194-01214,01216-01278,01280-01723,01928-01935,01940-01967,01972-02319,02324-02351,02356-02703,02708-02735,02740-03087,03092-03119,03124-03471,03476-03503,03512-03855,03860-03887,03892-04239,04244-04271,04280-07679]
PriorityJobFactor=10 PriorityTier=20 RootOnly=NO ReqResv=NO OverSubscribe=EXCLUSIVE
OverTimeLimit=NONE PreemptMode=OFF
State=UP TotalCPUs=253872 TotalNodes=7318 SelectTypeParameters=NONE
DefMemPerNode=UNLIMITED MaxMemPerNode=UNLIMITED

stud33@daint107:~> █
```

# Editing files

- vim or gvim (X version)
- emacs -nw or just emacs (X version)
- gedit (X only)



# Moving data to/from CSCS

- **scp**: Remote copy over SSH
  - Getting a file: `scp studXX@ela.cscs.ch:remotefile localfile`
  - Getting a directory: `scp -r studXX@ela.cscs.ch:remotedir localdir`
  - Sending a file: `scp localfile studXX@ela.cscs.ch:remotefile`
  - Sending a directory: `scp localdir studXX@ela.cscs.ch:remotedir`
- **rsync**: Synchronize files remotely over SSH
  - `rsync -avz studXX@ela.cscs.ch:remotedir/ localdir/`
  - `rsync -avz localdir/ studXX@ela.cscs.ch:remotedir/`
  - Pay attention to the slashes! *rsync behaves differently with or without slashes.*

# Summer school repository

All the material of the course is placed inside the following Github repo:

- <https://github.com/eth-cscs/SummerSchool2019>
- For instructions on how to clone and pull from the repository, check its front page.

Organization of the repository

- `miniapp/`: The different versions of the mini-app that you will work throughout the summer school + slides.
- `topics/`: The practical exercises of the different topics that will be covered during the the summer school + slides.
- `scripts/`: Useful scripts for the exercises and the mini-app.

Solutions:

- The solutions of the exercises and the mini-app will appear at the end of the summer school in subfolders named `solutions/` under each respective topic.